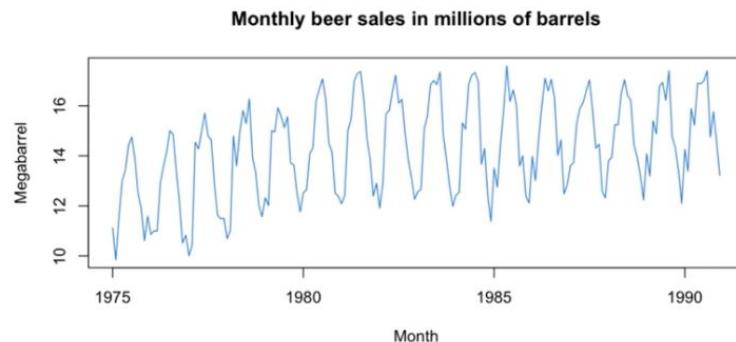


Time series Analysis

What is a time series

- Complex topic :
 - information theory
 - signal
 - probability/statistics
 - artificial intelligence/machine learning
 - optimization
 - Computer science (python)
 - Financial mathematics
- <http://www.laurentoudre.fr/ast.html>
- <http://www.laurentoudre.fr/signalmle.html>
- <https://scikit-learn.org/stable/>
- <https://github.com/fastai/course22>
- <https://github.com/fastai/fastbook>
- https://github.com/sduprey/initiation_python_finance
- https://github.com/sduprey/timeseries_ressources
- Continuous value (no sequence, discrete time data)
- Multivariate vector time series (brain signals ECG, multiple stocks)



Modeling vs Predicting : narrowing down the determinism

A framework which encompasses AI and econometry

$$y_t = f(y_{t-1}, \dots, y_{t-p}; \theta) + \varepsilon_t$$

- AI predilection domain: complex function with embeddings, features engineering and a small pure random part (see for example the Rossmann sales kaggle competition)
- Econometrics predilection domain : function very simple (just a scalar for a random walk) (try for example to predict the bitcoin price)

Modeling vs Predicting : narrowing down the determinism

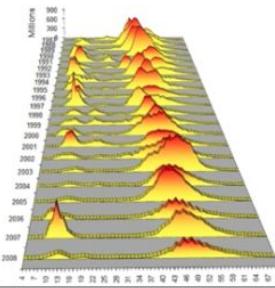
A framework which encompasses AI and econometry

$$y_t = f(y_{t-1}, \dots, y_{t-p}; \theta) + \varepsilon_t$$

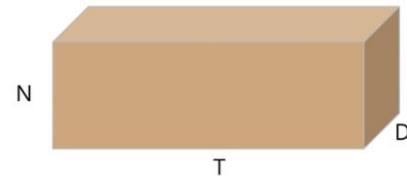
- Gives a functional form to your time series
- Gives you insight into the behavior of the time series
- “Why is the time series mean-reverting?”
- “Why does the time series grow unbounded?”
- “Is the time series predictable?”
- Someone who doesn’t know how to model might waste their time trying to predict something which is unpredictable! (e.g. a coin flip)

Dataframe (R, Pandas, PyTorch/TensorFlow tensors)

	New York City	London	Tokyo	Paris
1990-01-01	1	2	3	4
1990-01-02	5	6	7	8
1990-01-03	9	10	11	12
1990-01-04	13	14	15	16



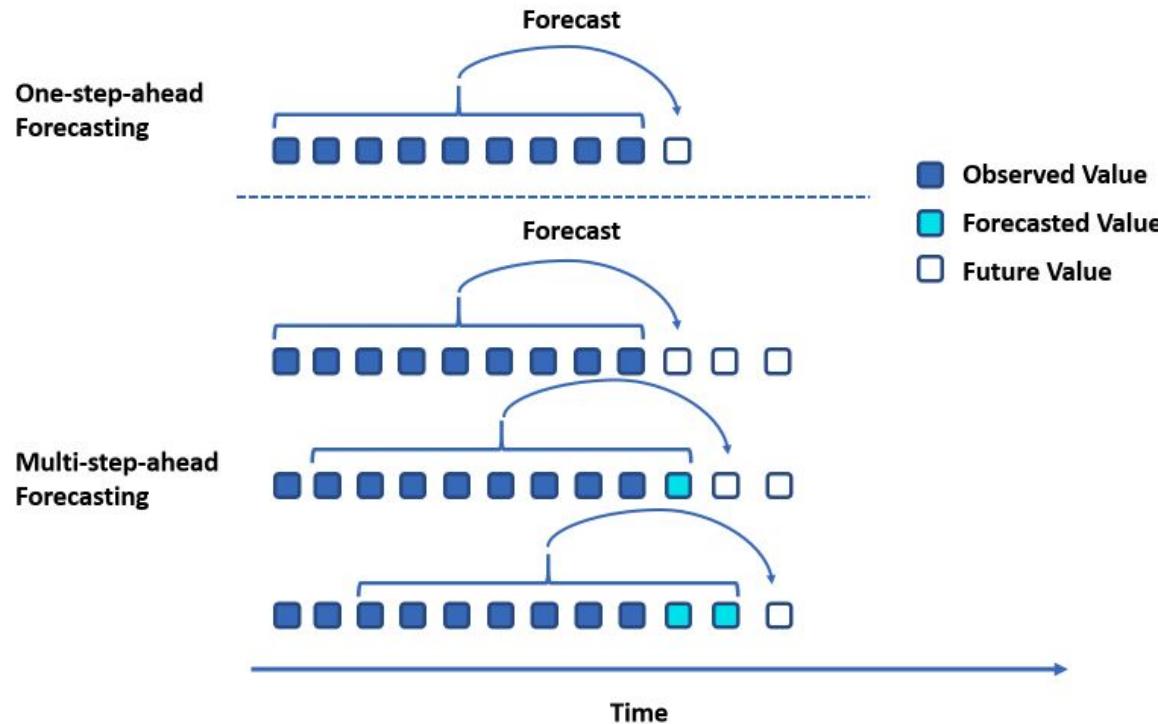
- $N \times T \times D$ (a box in 3-D space)
- Automatically think of this, whenever you see / hear “ $N \times T \times D$ ”
- It really helps, and makes it less abstract



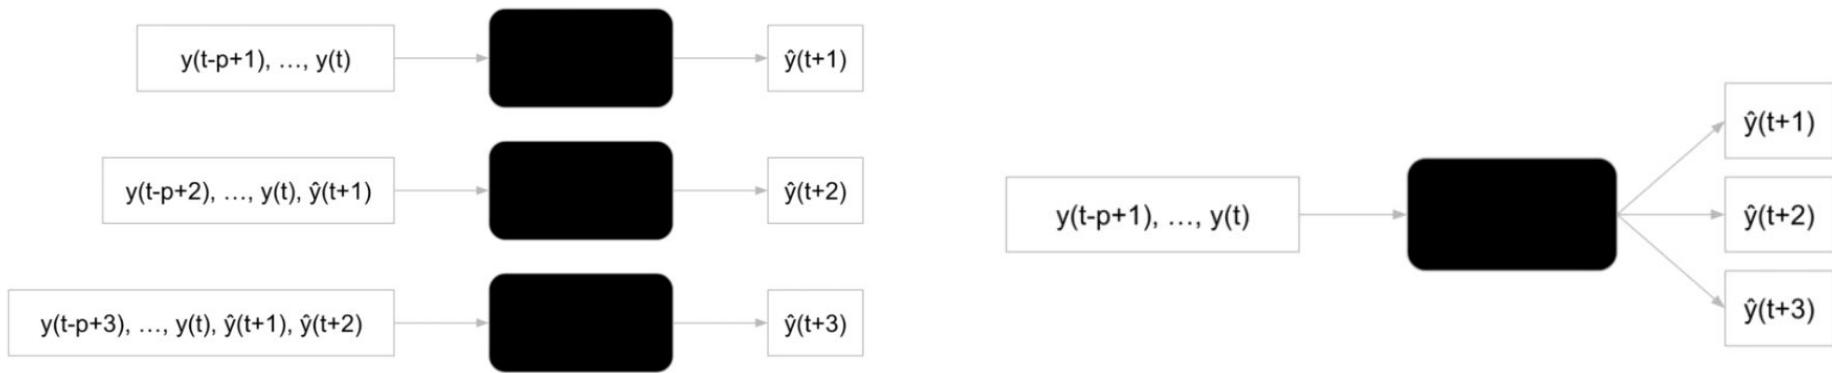
Generate random tensors/time series

- `np.random.randn(3,3,3)`
- Generate from different probability distributions
- Distribution fitting with scikit-learn

1-step forecast versus multi-steps forecast



Incremental multi-steps forward versus multi-output



Stochastic processes in a nutshell

- Stochastic processes are processes that proceed randomly in time.
- Rather than consider fixed random variables X , Y , etc. or even sequences of i.i.d random variables, we consider sequences X_0, X_1, X_2, \dots Where X_t represent some random quantity at time t .
- In general, the value X_t might depend on the quantity X_{t-1} at time $t-1$, or even the value X_s for other times $s < t$.
- Example: simple random walk .

Going more into maths:

- **Markov process**
- **Martingales**
- **Discretization of continuous processes**

Given an Itô process

$$dX(t) = \mu(t, X(t)) dt + \sigma(t, X(t)) dW(t),$$

and a time discretization $\{t_i \mid i = 0, \dots, n\}$ with $0 = t_0 < \dots < t_n$, then the time-discrete stochastic process \tilde{X} defined by

$$\tilde{X}(t_{i+1}) = \tilde{X}(t_i) + \mu(t_i, \tilde{X}(t_i)) \Delta t_i + \sigma(t_i, \tilde{X}(t_i)) \Delta W(t_i)$$

is called an *Euler-Maruyama scheme* of the process X (where $\Delta t_i := t_{i+1} - t_i$ and $\Delta W(t_i) := W(t_{i+1}) - W(t_i)$).

Random walk for instance

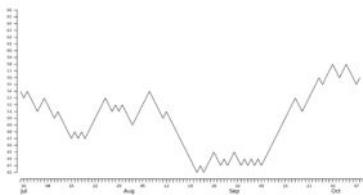
Discrete random walk

$p_0 = \text{some initial value}$

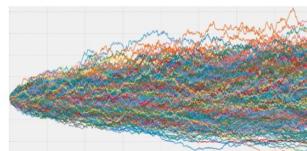
$p_1 = p_0 + e_1, \text{ where } e_1 \in \{-1, +1\}$

$p_2 = p_1 + e_2$

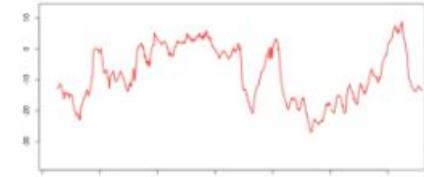
...



- Imagine yourself walking - you take one step **left** or **right** based on a coin flip - that's this random walk!
- Can't predict the future (50% chance of being correct)



Gaussian random walk



$p_0 = \text{some initial value}$

$p_1 = p_0 + e_1, \text{ where } e_i \sim \mathcal{N}(0, \sigma^2)$

$p_2 = p_1 + e_2$

...

Log Prices

- Consider a random walk with drift

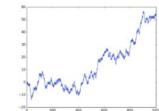
$$p_t = p_{t-1} + \mu + e_t, \quad e_t \sim \mathcal{N}(0, \sigma^2)$$

- Take $p(t-1)$ to the LHS - this is now the log return

$$r_t = p_t - p_{t-1} = \mu + e_t$$

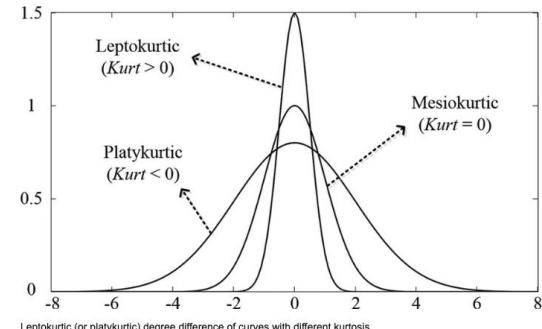
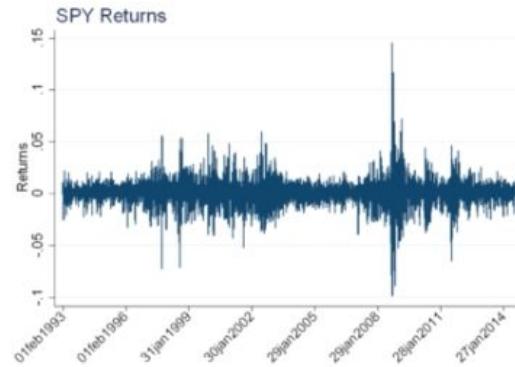
- The log return is therefore distributed as follows

$$r_t \sim \mathcal{N}(\mu, \sigma^2)$$



Stylized facts about financial time-series which invalidate the random walk hypothesis

- Volatility clustering
- Shock asymmetry
- Leptokurtic residuals



Markov property $p(w_t \mid w_{t-1}, w_{t-2}, \dots, w_0) = p(w_t \mid w_{t-1})$

Gaussian random walk

$$x(t) = x(t-1) + e(t), \quad e(t) \sim \mathcal{N}(0, \sigma^2)$$

$$x(t) \sim \mathcal{N}(x(t-1), \sigma^2)$$

$$\text{var}\{x(t+\tau)\} = ?$$

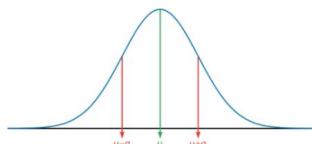
$$x(t+1) = x(t) + e(t+1)$$

$$x(t+2) = x(t+1) + e(t+2) = x(t) + e(t+1) + e(t+2)$$

...

$$x(t+\tau) = x(t) + e(t+1) + \dots + e(t+\tau)$$

Square root of time
growing confidence interval
Central limit theorem :
Converges to a gaussian

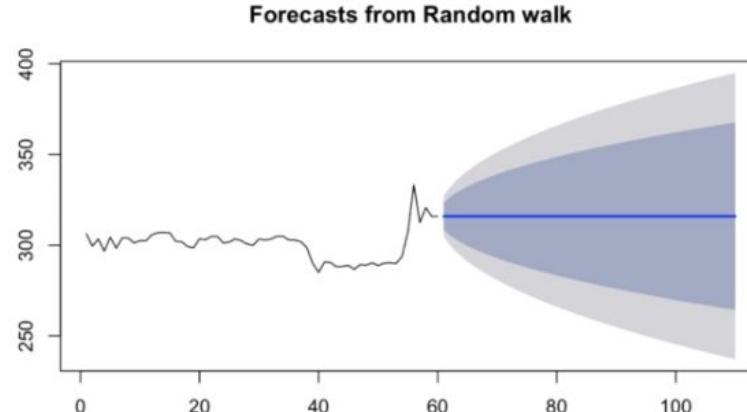
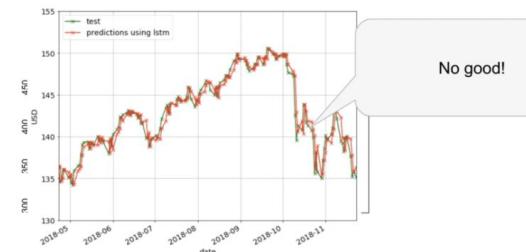


If $\text{var}\{e(t)\} = \sigma^2$, then $\text{var} \sum_{k=1}^{\tau} e(t+k) = \tau \sigma^2$, or $\text{sd}\{x(t+\tau)\} = \sqrt{\tau} \sigma$



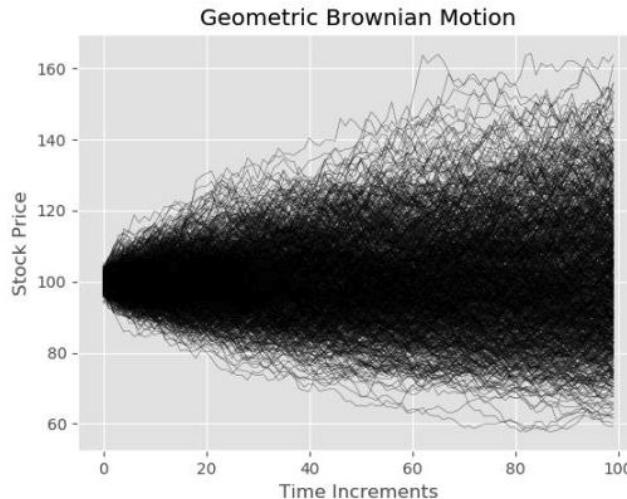
Naive forecast or the importance of a baseline

- Compare your results against existing state of the art
- Be careful :
 - the naïve forecast looks good for martingale processes
 - In-sample versus out-of-sample forecasts
 - You can beat the naive forecast in-sample but not out-sample
- **The naive forecast is the best for a random walk (for any martingales)**



Option pricing through Monte-Carlo simulation

- https://github.com/sduprey/timeseries_ressources/Black_Scholes_option_pricing_through_Monte_Carlo_simulation.ipynb



$$dS = S\mu dt + S\sigma dW(t)$$

$$dG = \left(\frac{\partial G}{\partial S} S\mu + \frac{\partial G}{\partial t} + \frac{1}{2} \frac{\partial^2 G}{\partial S^2} S^2 \sigma^2 \right) dt + \frac{\partial G}{\partial S} S\sigma dW(t)$$

$$\frac{\partial G}{\partial S} = \frac{1}{S} , \quad \frac{\partial G}{\partial t} = 0 , \quad \frac{\partial^2 G}{\partial S^2} = -\frac{1}{S^2}$$

$$\begin{aligned} dG &= \left(\frac{1}{S} S\mu + 0 - \frac{1}{2} \frac{1}{S^2} S^2 \sigma^2 \right) dt + \frac{1}{S} S\sigma dW(t) \\ &= (\mu - \frac{\sigma^2}{2}) dt + \sigma dW(t) \end{aligned}$$

$$lookback = e^{-rT} \left[\frac{1}{N} \sum_{i=1}^N (S_{max} - K)^+ \right]$$

Forecasting metrics (like regression)

- Sum of squared errors (max likelihood when the errors are normally distributed) $E = \sum_{i=1}^N (y_i - \hat{y}_i)^2$
- Mean squared error $E = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2 \approx E((y - \hat{y})^2)$
- Root mean squared error (same unit) $E = \sqrt{\frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2}$
- Mean absolute error $E = \frac{1}{N} \sum_{i=1}^N |y_i - \hat{y}_i|$

R squared

- Not an error - we want it to be bigger, not smaller
- If your model makes perfect predictions, then $MSE=0, R^2=1$
- $R^2=0$ means your model does no better than predicting \bar{y}

$$E = 1 - \frac{SSE}{SST} = 1 - \frac{MSE}{Var(Y)}$$

where $SST = \sum_{i=1}^N (y_i - \bar{y})^2$, $Var(Y) = \frac{1}{N} SST$

Scikit-learn

- For classification models: `model.score(X, Y)` returns accuracy
- For regression models: `model.score(X, Y)` returns R^2



Problems of relativity

- Accuracy is not the best metrics for an imbalanced classifier

$$E = \frac{1}{N} \sum_{i=1}^N \left| \frac{y_i - \hat{y}_i}{y_i} \right|$$

- MAPE (mean absolute percentage error)

$$E = \frac{1}{N} \sum_{i=1}^N \frac{|y_i - \hat{y}_i|}{(|y_i| + |\hat{y}_i|) / 2}$$

- sMAPE

Common transformations : analog to features engineering in artificial intelligence

- Power transform $y'(t) = y(t)^\gamma$
- Log transform $y'(t) = \log y(t)$ or $\log(y(t) + 1)$
- Box-Cox transform
$$y'(t) = \frac{y(t)^\lambda - 1}{\lambda} \quad \text{if } \lambda \neq 0$$
$$y'(t) = \log y(t) \quad \text{if } \lambda = 0$$

Since:

$$\lim_{\lambda \rightarrow 0} \frac{x^\lambda - 1}{\lambda} = \ln x$$

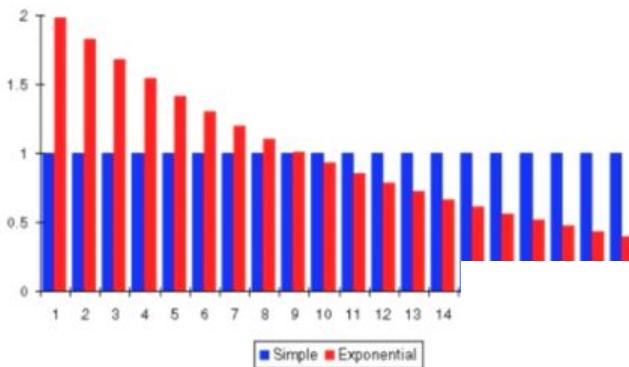
Why the log transform is fundamental

- Watch fast.ai on random forest : log uniformises large tail distributions
- In more standard econometrics, it is all about stationarity
- Rescaling data to more linear trends (decibels are the best example)
- It is part of a broader trend : features engineering for tabular data

https://github.com/fastai/fastbook/blob/master/09_tabular.ipynb

Simple/Exponentially weighted moving average

- The SMA is equally weighted, since for a window of size N, each point has a weight 1/N
- The EWMA gives each point a weight that decays exponentially



$$\text{Arithmetic mean : } \bar{x} = \frac{1}{T} \sum_{t=1}^T x_t$$

$$\text{EWMA : } \bar{x}_t = \alpha x_t + (1 - \alpha)\bar{x}_{t-1}, \text{ where } 0 \leq \alpha \leq 1$$

EWMA is tailoring the importance you give to more recent samples

We make $1/t$ which dwindle constant !

$$\bar{x}_t = \left(1 - \frac{1}{t}\right) \bar{x}_{t-1} + \frac{1}{t} x_t$$

$$\bar{x}_t = (1 - \alpha) \bar{x}_{t-1} + \alpha x_t$$

$$\begin{aligned}\bar{x}_t &= (1 - \alpha) \bar{x}_{t-1} + \alpha x_t \\ &= (1 - \alpha)[(1 - \alpha) \bar{x}_{t-2} + \alpha x_{t-1}] + \alpha x_t \\ &= (1 - \alpha)^2 \bar{x}_{t-2} + (1 - \alpha)\alpha x_{t-1} + \alpha x_t \\ &= (1 - \alpha)^2[(1 - \alpha) \bar{x}_{t-3} + \alpha x_{t-2}] + (1 - \alpha)\alpha x_{t-1} + \alpha x_t \\ &= (1 - \alpha)^3 \bar{x}_{t-3} + (1 - \alpha)^2 \alpha x_{t-2} + (1 - \alpha)\alpha x_{t-1} + \alpha x_t \\ &\dots \\ &= (1 - \alpha)^t \bar{x}_0 + \alpha \sum_{k=0}^{t-1} (1 - \alpha)^k x_{t-k}\end{aligned}$$

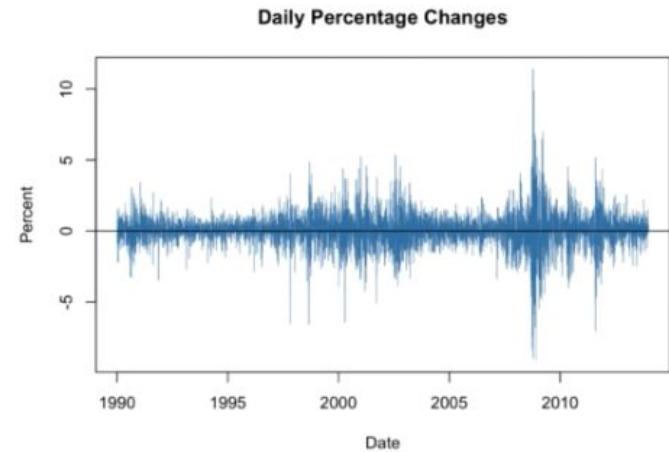
Why use moving average ?

Local assessment of mean/variance of
a time serie (volatility clustering)

```
# returns a Rolling object
rolling_window = df['GOOG'].rolling(window_size)

# returns a series/dataframe of rolling means
# can also calculate min, max, sum, var, etc.
rolling_window.mean()

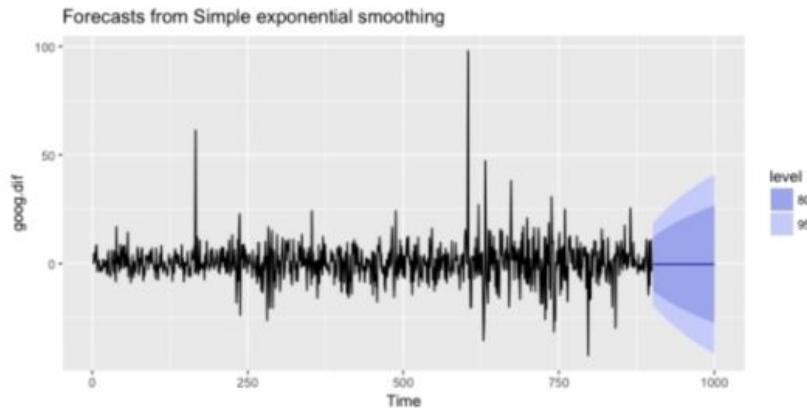
# multi-dimensional
covariance = df[['GOOG', 'AAPL']].rolling(50).cov()
```



Simple Exponential Smoothing

It makes the assumption that the process fluctuates around a mean value, which we assess through an EWMA (very efficient for martingales !)

$$\text{Level}(t+h) = \text{EWMA}(\text{Time series from } 1 \dots t)$$



$$\hat{y}_{t+1|t} = \alpha y_t + (1 - \alpha) \hat{y}_{t|t-1}$$

SES multi step forecast

Forecast Equation : $\hat{y}_{t+h|t} = l_t, h = 1, 2, 3\dots$

Smoothing Equation : $l_t = \alpha y_t + (1 - \alpha)l_{t-1}$

- Just the EWMA
- Notice: the original time indices are back!

```
from statsmodels.tsa.holtwinters import SimpleExpSmoothing  
  
# make the model - data is univariate  
ses = SimpleExpSmoothing(data)  
  
# 'fit' the model - returns a HoltWintersResults object  
result = ses.fit(smoothing_level=alpha, optimized=False)  
  
# in-sample prediction or out-of-sample forecast  
result.predict(start=start_date, end=end_date)
```

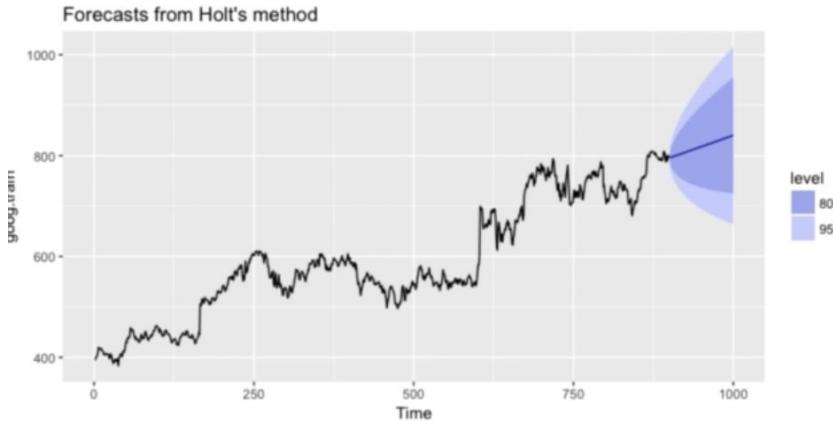
```
# get all in-sample predictions  
result.fittedvalues  
  
# forecast n steps ahead  
result.forecast(n)
```

If your data goes from $t=1\dots T$, then
the first forecast will be at time $T+1$

Holt's linear trend model

It makes the assumption that the process has a linear trend (level is the intercept, trend is the slope)

$$\text{Level}(t+h) = \text{EWMA}(\text{Level of time series from } 1 \dots t)$$
$$\text{Trend}(t+h) = \text{EWMA}(\text{Trend of time series from } 1 \dots t)$$



$$y = mx + b$$

$$y_t = \text{slope} \times t + y_0$$

Holt's linear trend model

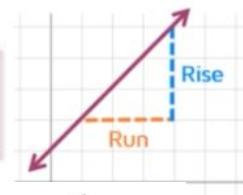
Forecast Equation : $\hat{y}_{t+h|t} = l_t + h b_t$

Level Equation : $l_t = \alpha y_t + (1 - \alpha)(l_{t-1} + b_{t-1})$

Trend Equation : $b_t = \beta(l_t - l_{t-1}) + (1 - \beta)b_{t-1}$

$$\text{slope in one step} = \frac{l_t - l_{t-1}}{1}$$

$$\text{Slope} = \frac{\text{Vertical Change}}{\text{Horizontal Change}} = \frac{\text{Rise}}{\text{Run}}$$

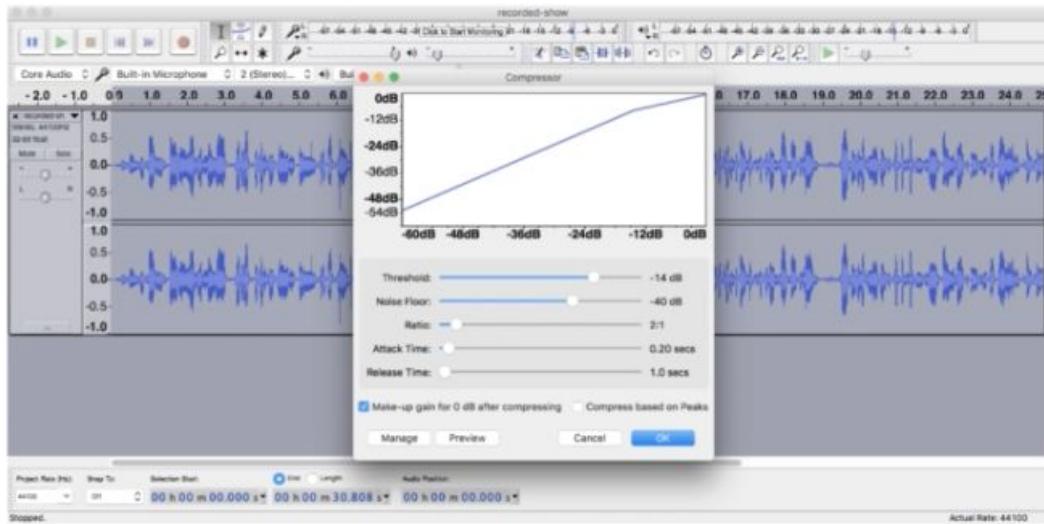


$$\text{one step increase due to trend} = l_{t-1} + 1 \times b_{t-1}$$

$$\text{Error} = \sum_{t=1}^T (y_t - \hat{y}_{t|t-1})^2$$

$$\alpha, \beta = \arg \min_{\alpha, \beta} \text{Error}$$

Alpha is a hyper parameter : like a frequency cut-off parameter in a signal filter

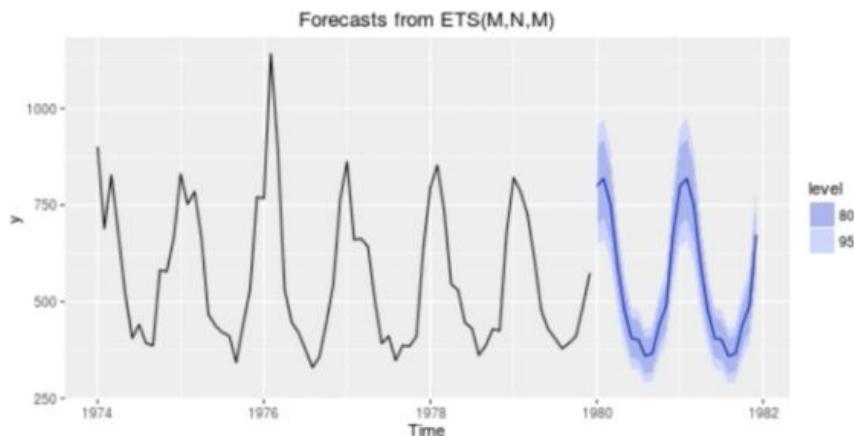


Holt-Winter model

$\text{Level}(t+h) = \text{EWMA}(\text{Level of time series from } 1 \dots t)$

$\text{Trend}(t+h) = \text{EWMA}(\text{Trend of time series from } 1 \dots t)$

$\text{Seasonal}(t+h) = \text{EWMA}(\text{Seasonal of time series from } 1 \dots t)$



- Additive: $y = \text{level} + \text{trend} + \text{seasonal}$
- Multiplicative $y = (\text{level} + \text{trend}) * \text{seasonal}$

Holt-Winter model (additif)

$$Forecast : \hat{y}_{t+h|t} = l_t + h b_t + s_{t+h-mk}$$

$$Level : \quad l_t = \alpha(y_t - s_{t-m}) + (1 - \alpha)(l_{t-1} + b_{t-1})$$

$$Trend : \quad b_t = \beta(l_t - l_{t-1}) + (1 - \beta)b_{t-1}$$

What is 'm'? The period of the cycle

$$Seasonality : \quad s_t = \gamma(y_t - l_{t-1} - b_{t-1}) + (1 - \gamma)s_{t-m}$$

What is 'k'?

$$\text{The formula: } k = \text{floor}\left(\frac{h-1}{m}\right) + 1$$

The intuition: finds the latest matching seasonal component (e.g. for March forecast, find the last-known March)

```
from statsmodels.tsa.holtwinters import ExponentialSmoothing

# make the model - data is univariate
# trend/seasonal args can be 'add' or 'mul'
model = ExponentialSmoothing(
    data, trend='add', seasonal='add', seasonal_periods=12)

# 'fit' the model - returns a HoltWintersResults object
result = model.fit()

# in-sample prediction or out-of-sample forecast
result.fittedvalues
result.forecast(n)
```

Explaining the k formula

$$\hat{y}_{t+h|t} = l_t + h b_t + s_{t+h-mk} \quad k = \text{floor}\left(\frac{h-1}{m}\right) + 1$$

- $t = 36, m = 12, h = 15$
- $k = \text{floor}[(15 - 1) / 12] + 1 = 2$
- Index for s: $t + h - mk = 36 + 15 - 12 * 2 = 27$

Month	Mar	...	Dec	Jan	...	Jan	Feb	Mar
t	27	...	36	37	...	49	50	51

Holt-Winter model (multiplicatif)

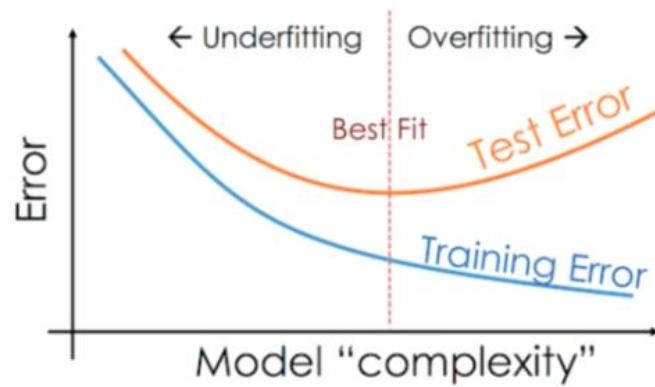
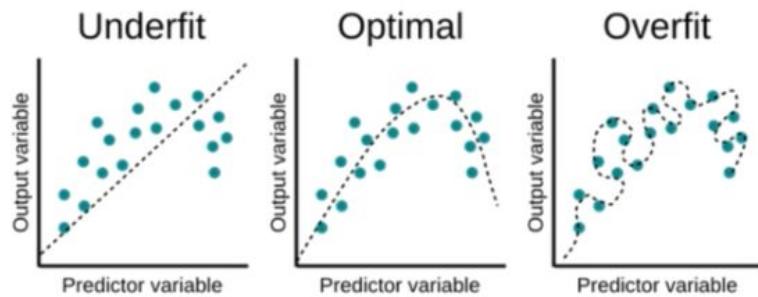
$$Forecast : \hat{y}_{t+h|t} = (l_t + hb_t)s_{t+h-mk}$$

$$Level : \quad l_t = \alpha \frac{y_t}{s_{t-m}} + (1 - \alpha)(l_{t-1} + b_{t-1})$$

$$Trend : \quad b_t = \beta(l_t - l_{t-1}) + (1 - \beta)b_{t-1}$$

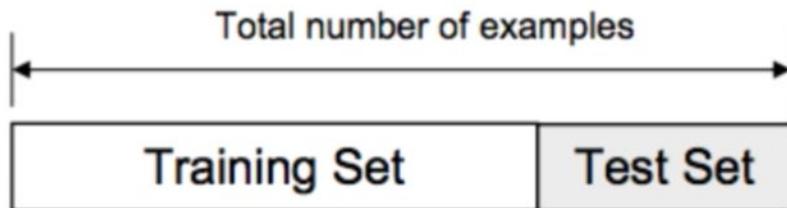
$$Seasonal : \quad s_t = \gamma \frac{y_t}{l_{t-1} + b_{t-1}} + (1 - \gamma)s_{t-m}$$

Walk-Forward validation



Single train-test splitting is wrong

- What happens when you try to optimize your model (via its configuration) by picking the best configuration for your test set?
- Your test set has effectively become “in-sample data”



Cross-validation to mitigate overfitting the validation set

- How about multiple validation sets? (Note: I'm conflating validation/test)
- Think of your true test set as real future data
- E.g. Private data in Kaggle contest, new customers to your website
- K times: train on K-1 parts, evaluate on the remaining part

5-fold CV		DATASET			
Estimation 1	Test	Train	Train	Train	Train
Estimation 2	Train	Test	Train	Train	Train
Estimation 3	Train	Train	Test	Train	Train
Estimation 4	Train	Train	Train	Test	Train
Estimation 5	Train	Train	Train	Train	Test

For time-series ?

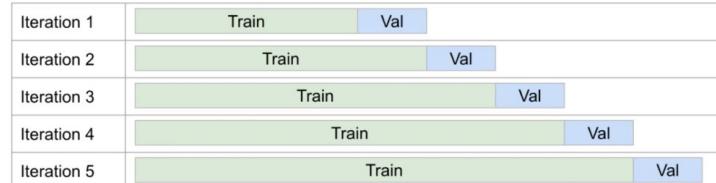
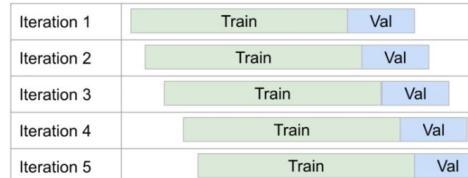
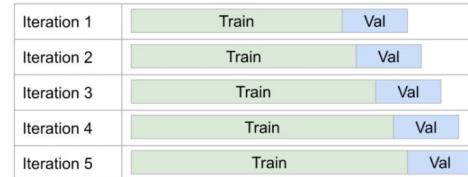
- Since the data is split randomly, you'll mix past and future data
- In the real-world, your model cannot be trained on future data!
- We can only use past data to predict future data

5-fold CV

	DATASET				
Estimation 1	Test	Train	Train	Train	Train
Estimation 2	Train	Test	Train	Train	Train
Estimation 3	Train	Train	Test	Train	Train
Estimation 4	Train	Train	Train	Test	Train
Estimation 5	Train	Train	Train	Train	Test

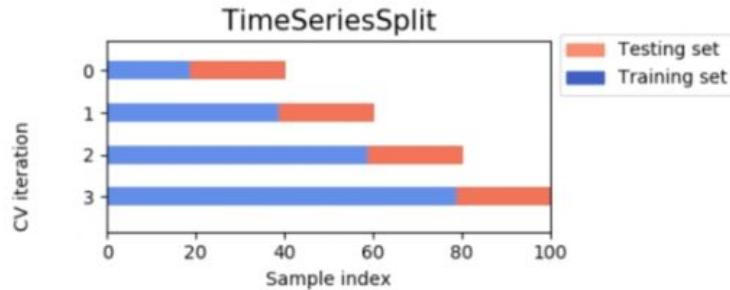


- Start with some minimum data to train on
- ... and validation forecast horizon $h \geq 1$
- Walk forward one step (train set becomes +1 bigger)
- Validate on the next h data points, etc etc. until you reach the end



TimeSeriesSplit (Scikit-Learn)

- Might save you from a **bit** of work, under the right circumstances
- To use with cross validation functions, your model must conform to scikit-learn interface (i.e. not statsmodels)
- Another limitation: not flexible - must use non-overlapping blocks
- All blocks are equal size - first train set is teeny tiny
- Do these reflect real-world operation?



Do your own code according to your needs and intuition !

We won't delve too much time using this approach into the different models we use : it complicates the script. We will just see an example script where you can plug your model afterwards.

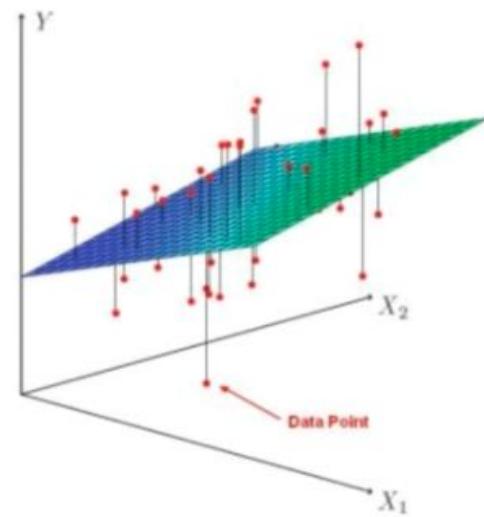
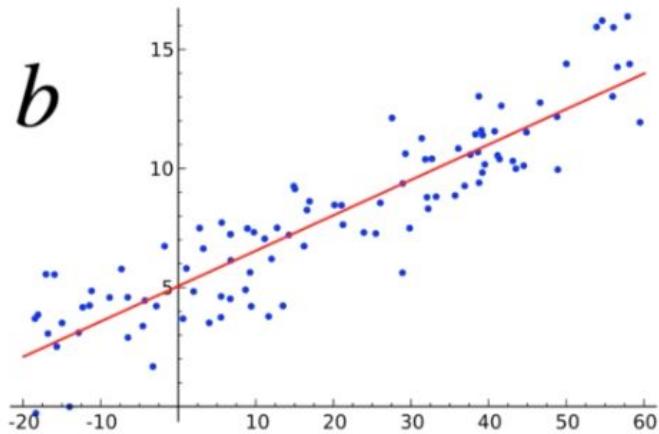
ARIMA vs Exponential smoothing

- Exponential smoothing is very specific (linear trends, seasonality)
- ARIMA imposes no such structure
- It is more “machine learning”-like

Autoregressive models

$$\hat{y} = w_1x_1 + w_2x_2 + b$$

$$\hat{y} = mx + b$$



AR(p) : auto-regressive processus

$$\hat{y}_t = b + \varphi_1 y_{t-1} + \varphi_2 y_{t-2} + \dots + \varphi_p y_{t-p}$$

- Suppose our data is: y_1, y_2, \dots, y_{10}
- In ML, we say X has shape N x D, but for ARIMA we'll stick with D == p

y1	y2	y3
y2	y3	y4
y3	y4	y5
y4	y5	y6
y5	y6	y7
y6	y7	y8
y7	y8	y9

This is our "X"

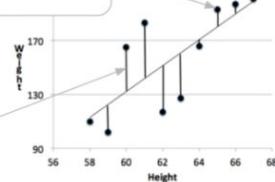
y4
y5
y6
y7
y8
y9
y10

This is our "Y"

$$y_t = b + \varphi_1 y_{t-1} + \varphi_2 y_{t-2} + \dots + \varphi_p y_{t-p} + \varepsilon_t$$
$$\varepsilon_t \sim \mathcal{N}(0, \sigma^2)$$
$$\hat{y}_t = E(y_t)$$

y (no hat) is the true data point

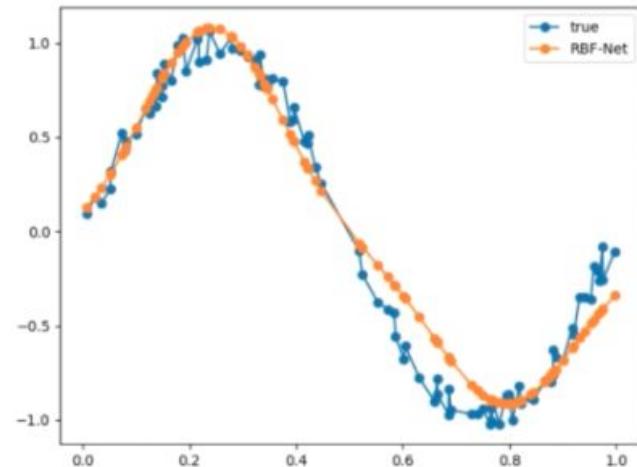
These are the errors



Machine Learning is the next step

- Linear models aren't that powerful (only lines or planes)
- Why stick to linear regression?
- ARIMA helps us understand the modeling and statistical properties

```
model = NeuralNetwork()  
model.fit(X, Y)  
  
model = RandomForest()  
model.fit(X, Y)
```



MA(q)

$$y_t = c + \varepsilon_t + \theta_1 \varepsilon_{t-1} + \theta_2 \varepsilon_{t-2} + \dots + \theta_q \varepsilon_{t-q}$$

Moving around the average c:

.

The diagram illustrates the decomposition of the expected value of y_t . A box labeled "Zero!" has arrows pointing to each term in the equation $E(y_t) = E(c + \varepsilon_t + \theta_1 \varepsilon_{t-1} + \theta_2 \varepsilon_{t-2} + \dots + \theta_q \varepsilon_{t-q}) = c$. The term c is also explicitly labeled at the end of the equation.

$$\begin{aligned} E(y_t) &= E(c + \varepsilon_t + \theta_1 \varepsilon_{t-1} + \theta_2 \varepsilon_{t-2} + \dots + \theta_q \varepsilon_{t-q}) \\ &= c \end{aligned}$$

ARMA(p,q)

- ARMA(p, q) = AR(p) + MA(q)

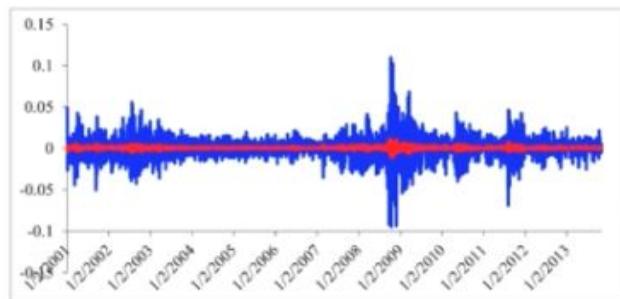
$$y_t = b + \boxed{\varphi_1 y_{t-1} + \dots + \varphi_p y_{t-p}} + \boxed{\theta_1 \varepsilon_{t-1} + \dots + \theta_q \varepsilon_{t-q}} + \varepsilon_t$$

The diagram illustrates the decomposition of an ARMA(p,q) process into its AR(p) and MA(q) components. Two arrows originate from the text 'ARMA(p, q) = AR(p) + MA(q)' and point to the two boxed terms in the equation below. The first arrow points to the term $\varphi_1 y_{t-1} + \dots + \varphi_p y_{t-p}$, which represents the AR(p) component. The second arrow points to the term $\theta_1 \varepsilon_{t-1} + \dots + \theta_q \varepsilon_{t-q}$, which represents the MA(q) component.

Differencing : detrending

- Where have we seen this?

- Log returns - the difference of log prices: $r_t = p_t - p_{t-1}$
- Holt's Linear Trend Model: $b_t = \beta(l_t - l_{t-1}) + (1-\beta)b_{t-1}$

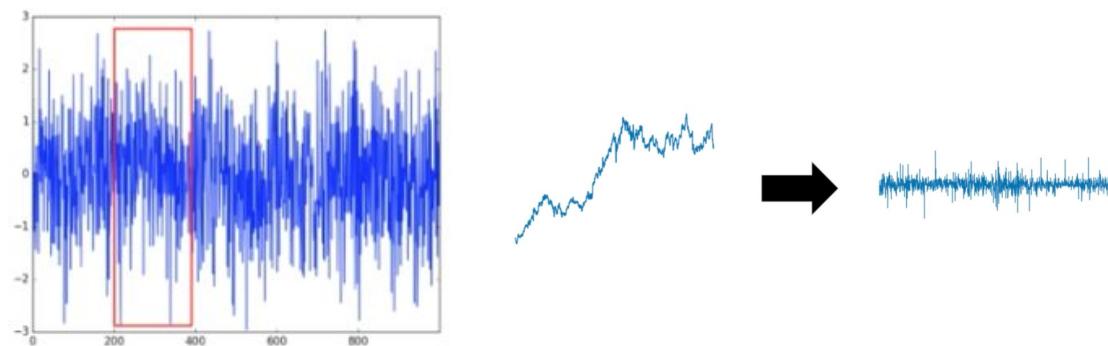


Given : $\{y_t\} = \{y_1, y_2, \dots, y_T\}$ (some time series)

Differenced Series : $\Delta y_t = y_t - y_{t-1}$

ARMA modelling needs stationary time series

- When fitting an ARMA model, we want the data to be close to stationary
- Stationary = Does not change over time
- Stationarity is nice: mean, variance, autocorrelation, ... will be constant over time
 - Recall: linear models fit well when there is strong correlation between inputs / output
- Each “window” of the time series is like a “training point” for fitting the model



I(d) and ARIMA(p,d,q)

- An I(d) process is a process that is stationary after differencing d times
- We say it's integrated to order d
- ARIMA(p, d, q) is just a model where we've differenced d times before applying ARMA(p, q)

Random Walk

- ARIMA(p, 0, 0) is AR(p)
- It's also ARMA(p, 0)
- ARIMA(0, 0, q) is ARMA(0, q) and MA(q)
- ARIMA(0, d, 0) is I(d)

- ARIMA(0, 1, 0) is I(1) and this is a random walk

Differenced time series

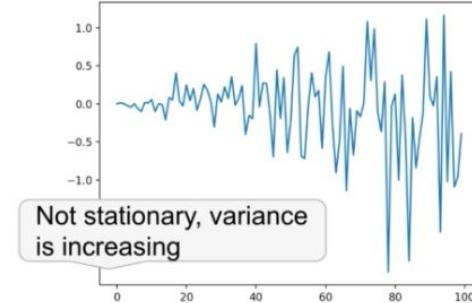
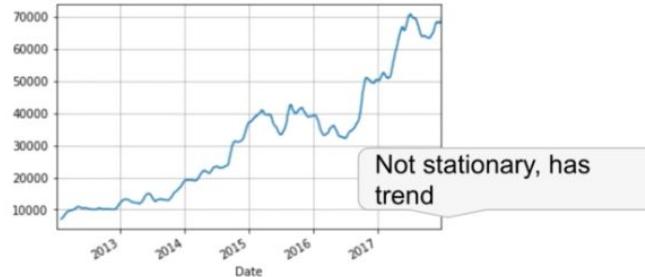
Absence of AR and MA components

$$\Delta y_t = \varepsilon_t$$

$$y_t - y_{t-1} = \varepsilon_t$$

$$y_t = y_{t-1} + \varepsilon_t$$

Stationarity

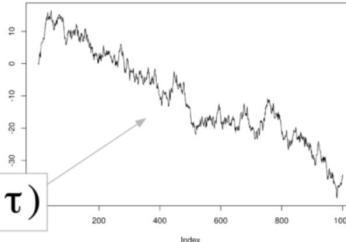


- Loosely, the distribution of the random variables in the time series does not change over time
- E.g. mean and variance will always be the same

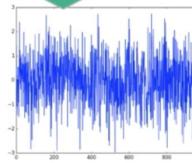
- If a time series is nonstationarity, then you might need different models at different points in time!
- For nonstationary time series, you can't treat data points at different times like "samples" (e.g. computing the mean, variance)

You can't compute "the mean" from the data because the mean is changing!

$$\mu_Y(t) \neq \mu_Y(t + \tau)$$



Here computing the mean across time is OK!

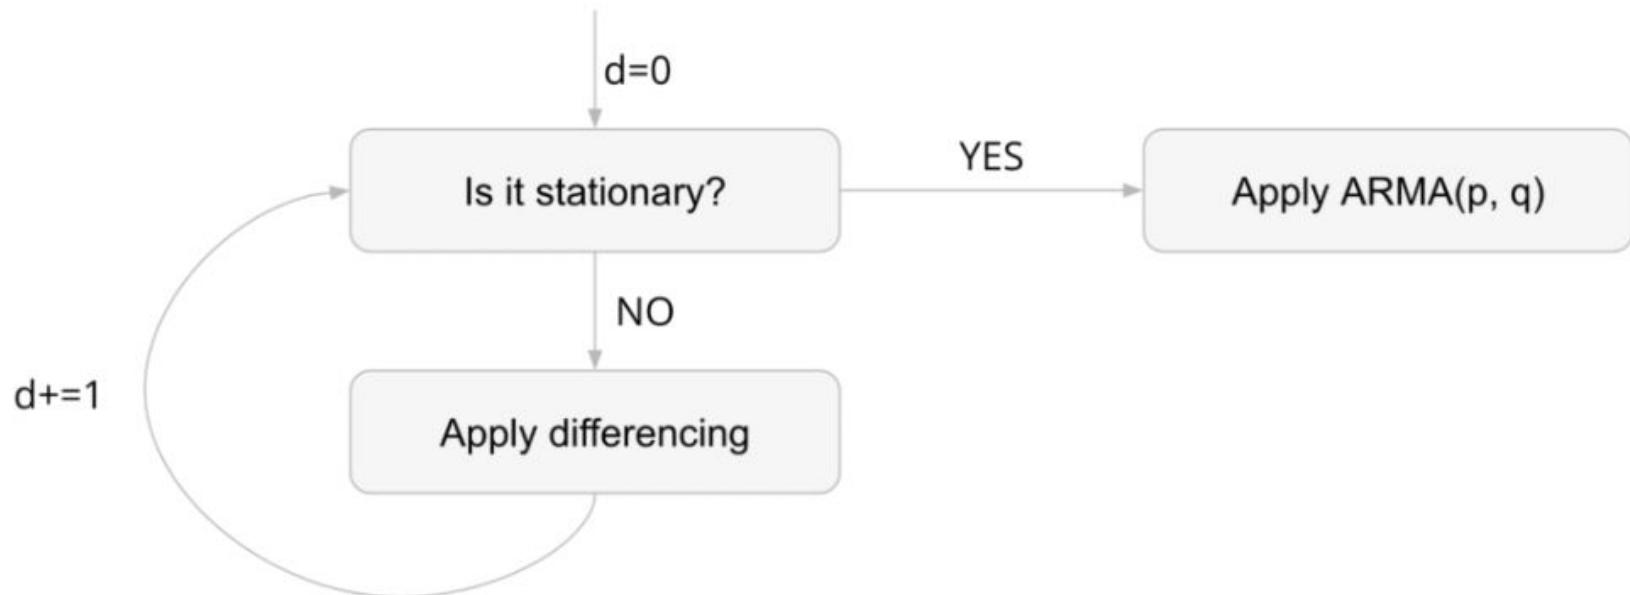


Testing for stationarity

- We use the Augmented Dickey-Fuller Test (ADF Test)
- Think of it like an API:
 - Given: null hypothesis, alternative hypothesis
 - Input: time series, Output: p-value
 - Action: accept or reject the null hypothesis
- For ADF test:
 - Null: time series is non-stationary
 - Alternative: time series is stationary



How to use ADF test in selecting d in ARIMA



Strong vs Weak

- Strong: the entire distribution does not change over time

$$F_Y(y_{t_1+\tau}, y_{t_2+\tau}, \dots, y_{t_n+\tau}) = F_Y(y_{t_1}, y_{t_2}, \dots, y_{t_n}), \forall \tau, t_1, t_2, \dots, t_n$$

- Weak : First order (mean) and second-order statistics (covariance) stay the same

- The mean does not change over time

$$\mu_Y(t) = \mu_Y(t + \tau) \text{ for all } \tau$$

- The autocovariance does not change over time

$$K_{YY}(t_1, t_2) = K_{YY}(t_1 - t_2, 0) \text{ for all } t_1, t_2$$

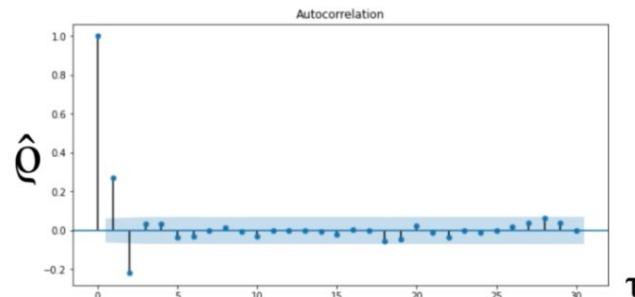
Autocovariance : $\text{cov}(Y_{t_1}, Y_{t_2})$

Autocorrelation function

Autocorrelation : $\frac{\text{cov}(Y_{t_1}, Y_{t_2})}{\sigma_Y(t_1)\sigma_Y(t_2)}$

Stationary : $\rho(Y(t_1), Y(t_2)) = \rho(t_1 - t_2) = \rho(\tau)$

- Also known as correlogram
- Autocorrelation is to autocovariance as correlation is to covariance
- Auto = Self (both RVs come from the same time series)



How to determine q in MA(q)

- Assign q to be the maximum non-zero lag
- E.g. in below chart, q = 2
- Usually, the ACF for lags < q are also non-zero

- This can be derived mathematically! (we won't do it right now)
- For MA(1):

$$y_t = c + \theta_1 \varepsilon_{t-1} + \varepsilon_t, \text{ where } \varepsilon_t \sim N(0, \sigma^2)$$

Why does it work ?

$$\varrho(1) = \frac{\theta_1}{1+\theta_1^2}, \quad \varrho(\tau) = 0 \text{ for } \tau > 1$$

- For MA(2):

$$\varrho(1) = \frac{\theta_1 + \theta_1 \theta_2}{1+\theta_1^2+\theta_2^2}, \quad \varrho(2) = \frac{\theta_2}{1+\theta_1^2+\theta_2^2}, \quad \varrho(\tau) = 0 \text{ for } \tau > 2$$

PACF of order p of AR(p)

- Definition: The PACF at lag τ is the autocorrelation between $Y(t)$ and $Y(t+\tau)$, *conditioned on* $Y(t+1), Y(t+2), \dots, Y(t+\tau-1)$



$$\varphi(\tau, \tau) = \text{corr}(Y_{t+\tau} - \hat{Y}_{t+\tau}, Y_t - \hat{Y}_t)$$

$$\hat{Y}_{t+\tau} = \beta_0 + \beta_1 Y_{t+1} + \beta_2 Y_{t+2} + \dots + \beta_{\tau-1} Y_{t+\tau-1}$$

$$\hat{Y}_t = \beta_0' + \beta_1' Y_{t+1} + \beta_2' Y_{t+2} + \dots + \beta_{\tau-1}' Y_{t+\tau-1}$$

GARCH theory

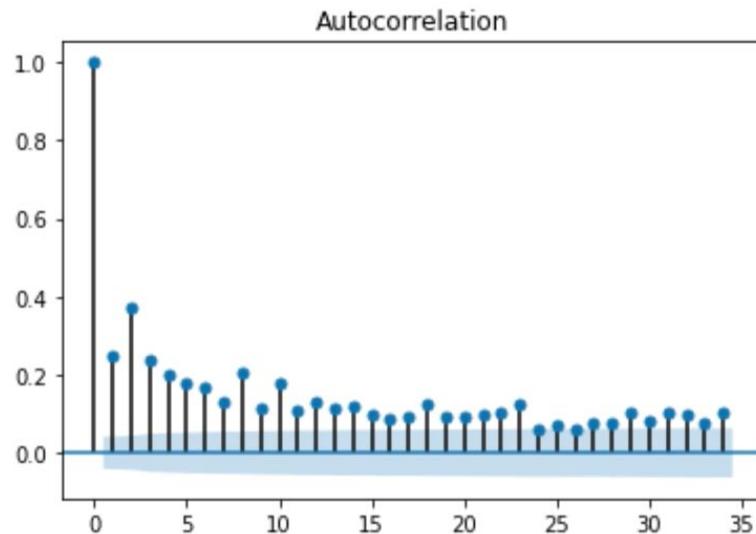
ARIMA: used to model
time series mean
(variance = 0)

$$y_t = f(y_{t-1}, \dots, y_{t-p}; \theta) + \varepsilon_t$$

GARCH: used to model
time series variance
(mean = 0)

ACF of squared returns show autoregressivity in variance

- For stocks, log returns ACF shows randomness ($I(1)$ is the best model)
- The ACF of squared log returns does not look random at all



ARCH(1)

ARCH(1)

Could be $N(0, 1)$, but not necessary

$$E[z_t] = 0, E[z_t^2] = 1$$

Time series we want to model

$$\varepsilon_t = z_t \sqrt{\omega + \alpha_1 \varepsilon_{t-1}^2}$$

Bias term

Autoregressive coefficient

$$\varepsilon_t = z_t \sigma_t \Rightarrow \text{epsilon has variance sigma squared}$$

$$\varepsilon_t^2 = z_t^2 (\omega + \alpha_1 \varepsilon_{t-1}^2)$$

$$\varepsilon_t = z_t \sigma_t$$

$$\frac{\varepsilon_t^2}{z_t^2} = \sigma_t^2 = \omega + \alpha_1 \varepsilon_{t-1}^2$$

- There must be constraints, since variance cannot be negative

$$\varepsilon_t = z_t \sqrt{\omega + \alpha_1 \varepsilon_{t-1}^2}$$

- General constraints

$$\omega > 0, \quad 0 \leq \alpha_1 \leq 1$$

- Constraint for stationarity and finite variance

$$\alpha_1 < 1$$

$$E \left[\varepsilon_t^2 \mid \varepsilon_{t-1}, \varepsilon_{t-2}, \dots, \varepsilon_{t-p} \right] = \sigma_t^2$$

$$E \left[\varepsilon_t^2 \right] = ? \text{ (call it } \sigma^2)$$

$$E \left[\varepsilon_t^2 \right] = E \left[z_t^2 \right] E \left[\omega + \alpha_1 \varepsilon_{t-1}^2 \right]$$

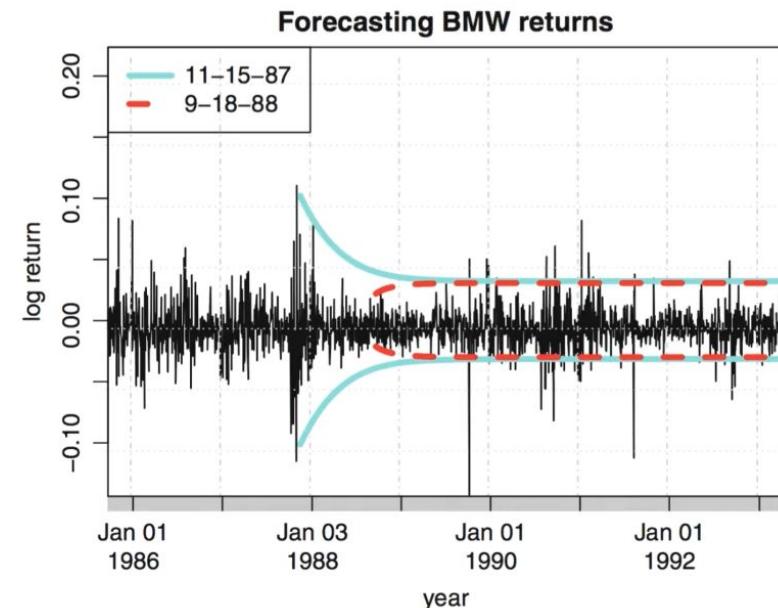
$$\sigma^2 = 1 \times (\omega + \alpha_1 \sigma^2)$$

$$(1 - \alpha_1) \sigma^2 = \omega$$

$$\sigma^2 = \frac{\omega}{1 - \alpha_1}$$

Long term forecast converge to the unconditional variance

- Our forecast of variance will also converge to unconditional variance
- Meaning: ARCH and GARCH are useful for short-term forecasts
- Makes sense: we cannot predict
Elon Musk making a tweet about
Bitcoin, or some country banning
a cryptocurrency exchange
- Obvious example: COVID-19
- Imagine if time series models could
predict world-scale events!
(of course, this is unreasonable)
- Aside from insider trading, you
probably cannot predict these



ARCH(p)

$$\varepsilon_t = z_t \sqrt{\omega + \sum_{\tau=1}^p \alpha_\tau \varepsilon_{t-\tau}^2}$$

$$E \left[\varepsilon_t \mid \varepsilon_{t-1}, \varepsilon_{t-2}, \dots, \varepsilon_{t-p} \right] = 0$$

$$E \left[\varepsilon_t \right] = 0$$

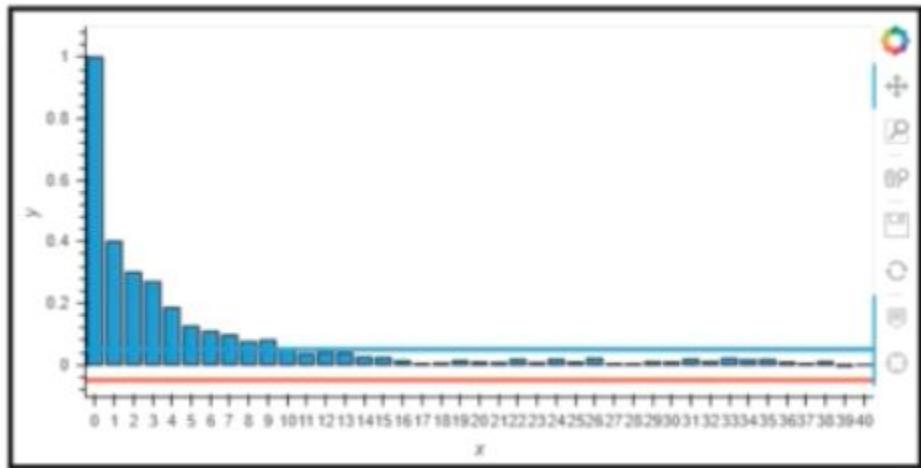
$$cov(\varepsilon_t, \varepsilon_{t-s}) = 0, \quad \forall s \neq 0$$

$$\varepsilon_t \perp \varepsilon_{t-s} \rightarrow cov(\varepsilon_t, \varepsilon_{t-s}) = 0$$

$$cov(\varepsilon_t, \varepsilon_{t-s}) = 0 \not\rightarrow \varepsilon_t \perp \varepsilon_{t-s}$$

ACF of squared epsilons : geometric decay

$$\rho_{\varepsilon^2}(h) = \alpha_1^{|h|}$$



- Unconditional mean is zero
- Conditional mean is zero
- ACF of series shows no correlation with lags (just like stock returns)
- ACF of squared series does show correlation (just like stock returns)
- Allows us to model volatility clustering

GARCH(p,q) : persistent volatility, less HFT moves

ARCH(p) : AR(p)

GARCH(p,q) : ARMA(p,q)

There are other ways to form this analogy, but I think you get the idea...

$$\sigma_t^2 = \omega + \sum_{k=1}^p \alpha_k \varepsilon_{t-k}^2 + \sum_{k=1}^q \beta_k \sigma_{t-k}^2$$

$$\varepsilon_t = z_t \sigma_t$$

$$\sigma_t^2 = \omega + \alpha_1 \varepsilon_{t-1}^2 + \beta_1 \sigma_{t-1}^2$$

- Consider a GARCH(1,1) for simplicity (actually, this is the most common)
- ε_{t-1} term is random, σ_{t-1} term is not random
- The “randomness” comes from z_t
- Imagine $\beta_1 = 0.9$ - then σ_t would decay slowly over time

GARCH(1,1) as ARMA(1,1)

$$\sigma_t^2 = \omega + \alpha_1 \varepsilon_{t-1}^2 + \beta_1 \sigma_{t-1}^2 \quad \eta_t = \varepsilon_t^2 - \sigma_t^2$$

$$\sigma_t^2 = \omega + \alpha_1 \varepsilon_{t-1}^2 + \beta_1 (\varepsilon_{t-1}^2 - \eta_{t-1})$$

$$\sigma_t^2 = \omega + (\alpha_1 + \beta_1) \varepsilon_{t-1}^2 - \beta_1 \eta_{t-1}$$

$$\sigma_t^2 + \eta_t = \omega + (\alpha_1 + \beta_1) \varepsilon_{t-1}^2 - \beta_1 \eta_{t-1} + \eta_t$$

$$\varepsilon_t^2 = \omega + (\alpha_1 + \beta_1) \varepsilon_{t-1}^2 - \beta_1 \eta_{t-1} + \eta_t$$

$$\eta_t = \varepsilon_t^2 - \sigma_t^2$$

$$\sigma_t^2 = \omega + \alpha_1 \varepsilon_{t-1}^2$$

$$\sigma_t^2 + \eta_t = \omega + \alpha_1 \varepsilon_{t-1}^2 + \eta_t$$

$$\varepsilon_t^2 = \omega + \alpha_1 \varepsilon_{t-1}^2 + \eta_t$$

$$y_t^2 = \omega + \alpha_1 y_{t-1}^2 + \eta_t, \text{ where } y_t = \varepsilon_t$$

A deep learning approach to GARCH

- 1) How does it make predictions? (i.e. go from input to output)
Usually some equation involving model parameters (ω, α, β for GARCH)
- 2) How do we find those model parameters?
E.g. Deep Learning: “call fit” / “gradient descent”

- For deep learning, ARIMA, and many other ML models, the objective is based on the log-likelihood
- Maximum Likelihood Estimation (MLE): maximize the likelihood wrt model parameters - this is the “training process”
- Regression: typically minimize $\sum_i (y_i - \hat{y}_i)^2$ (squared error)
- Equivalent to maximizing likelihood when errors are normal
- This requires another assumption! The variance of errors is constant

Maximum likelihood in case of homoscedasticity

Assumption : $y_t \sim \mathcal{N}(\hat{y}_t, \sigma^2)$

Equivalent : $y_t = \hat{y}_t + \varepsilon_t$, where $\varepsilon_t \sim \mathcal{N}(0, \sigma^2)$

$$\hat{y}_t(\theta) = f(y_{t-1}, \dots, y_{t-p}; \theta)$$

$$L(\theta) = \prod_{t=1}^T \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{1}{2} \frac{(y_t - \hat{y}_t)^2}{\sigma^2}\right)$$

$$l(\theta) = \log L(\theta) = \sum_{t=1}^T \left\{ -\frac{1}{2} \log(2\pi\sigma^2) - \frac{1}{2} \frac{(y_t - \hat{y}_t)^2}{\sigma^2} \right\}$$

$$\theta^* = \arg \max_{\theta} \sum_{t=1}^T \left\{ -\frac{1}{2} \log(2\pi\sigma^2) - \frac{1}{2} \frac{(y_t - \hat{y}_t)^2}{\sigma^2} \right\}$$

$$\theta^* = \arg \max_{\theta} \sum_{t=1}^T \left\{ -\frac{1}{2} \frac{(y_t - \hat{y}_t)^2}{\sigma^2} \right\}$$

$$\theta^* = \arg \max_{\theta} \sum_{t=1}^T \left\{ -(y_t - \hat{y}_t)^2 \right\}$$

$$\theta^* = \arg \min_{\theta} \sum_{t=1}^T \left\{ (y_t - \hat{y}_t)^2 \right\}$$

Max likelihood in case of hetero scedasticity

$$\theta^* = \arg \max_{\theta} \sum_{t=1}^T \left\{ -\frac{1}{2} \log(2\pi\sigma_t^2) - \frac{1}{2} \frac{(y_t - \hat{y}_t)^2}{\sigma_t^2} \right\}$$

$$\theta^* = \arg \max_{\theta} \sum_{t=1}^T \left\{ -\frac{1}{2} \log(\sigma_t^2) - \frac{1}{2} \frac{(y_t - \hat{y}_t)^2}{\sigma_t^2} \right\}$$

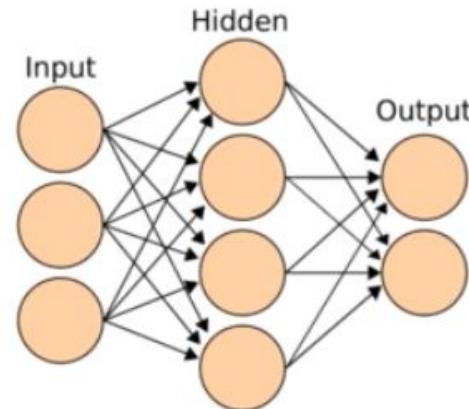
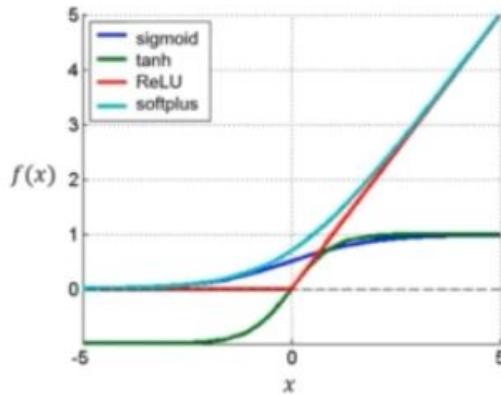
$$\theta^* = \arg \min_{\theta} \sum_{t=1}^T \left\{ \log(\sigma_t^2) + \frac{(y_t - \hat{y}_t)^2}{\sigma_t^2} \right\}$$

$$\theta^* = \arg \min_{\theta} \sum_{t=1}^T \left\{ \log(\sigma_t^2) + \frac{\varepsilon_t^2}{\sigma_t^2} \right\}$$

- Can plug into any stock optimizer (e.g. L-BFGS) - see Scipy docs
- We also need to include constraints on ω, α, β
- Easy to plug in with any optimizer
- Note: Derivation based on normal, but similar steps can be done for t-distribution, skewed t-distribution, etc.

The deep learning way

- Why stick with GARCH? (a linear model)
- We can parameterize σ_t using a neural network
- Just need to know how to make custom loss functions in TF2/PyTorch/...
- How to constraint output to be positive?
- Use positive-only activation (e.g. softplus)



Prerequisites for machine learning

Tensorflow 2.0



Tensorflow 2.0: Deep Learning and Artificial Intelligence

Lazy Programmer Inc., Lazy Programmer Team

PyTorch



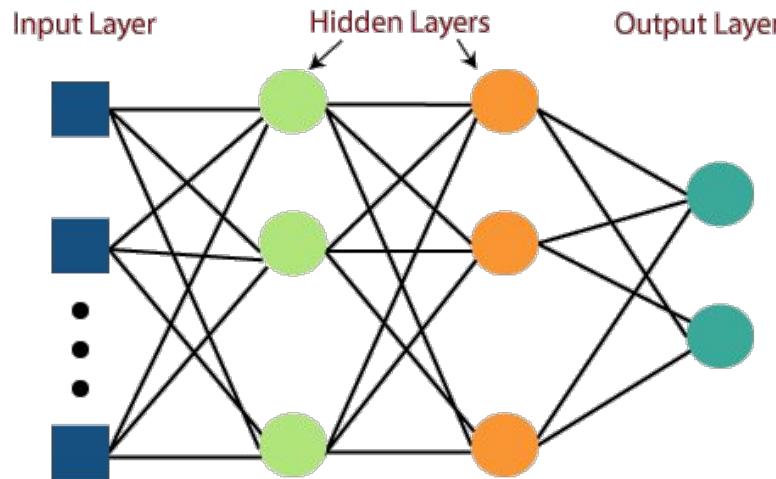
PyTorch: Deep Learning and Artificial Intelligence

Lazy Programmer Team, Lazy Programmer Inc.

In-depth deep learning series (up to CNNs and RNNs)



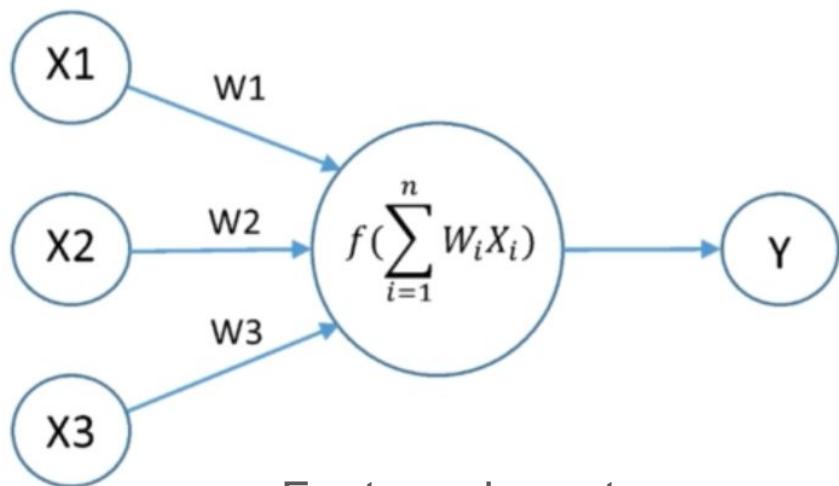
Multi-Layer perceptron



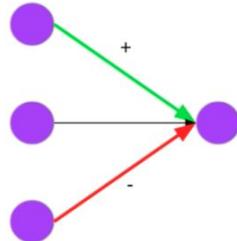
- <https://github.com/fastai/course22/blob/master/04-how-does-a-neural-net-really-work.ipynb>
- <https://developers.google.com/machine-learning/crash-course/introduction-to-neural-networks/video-lecture?hl=fr>

A line : $ax + b$

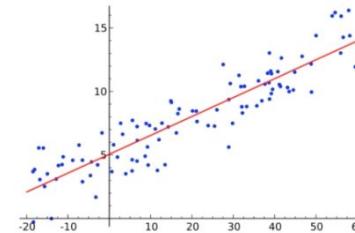
A neuron : $\sigma(w^T x + b)$



Features importance



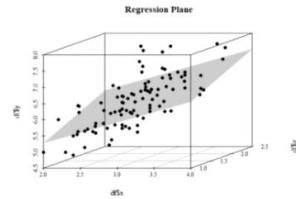
1d regression



$$\hat{y} = mx + b \text{ (or } ax + b\text{)}$$

2d regression

$$\hat{y} = w_1 x_1 + w_2 x_2 + b$$

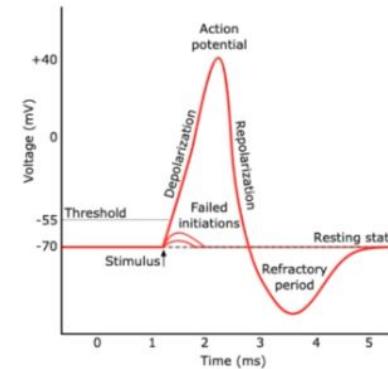
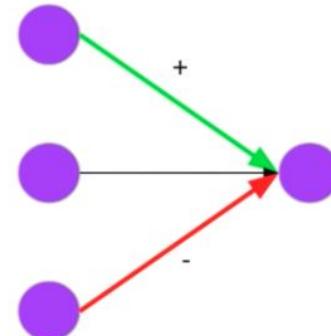


On the importance of introducing non-linearities through activation

- <https://github.com/fastai/course22/blob/master/04-how-does-a-neural-net-really-work.ipynb>
- [Introduction aux réseaux de neurones](#) (google)

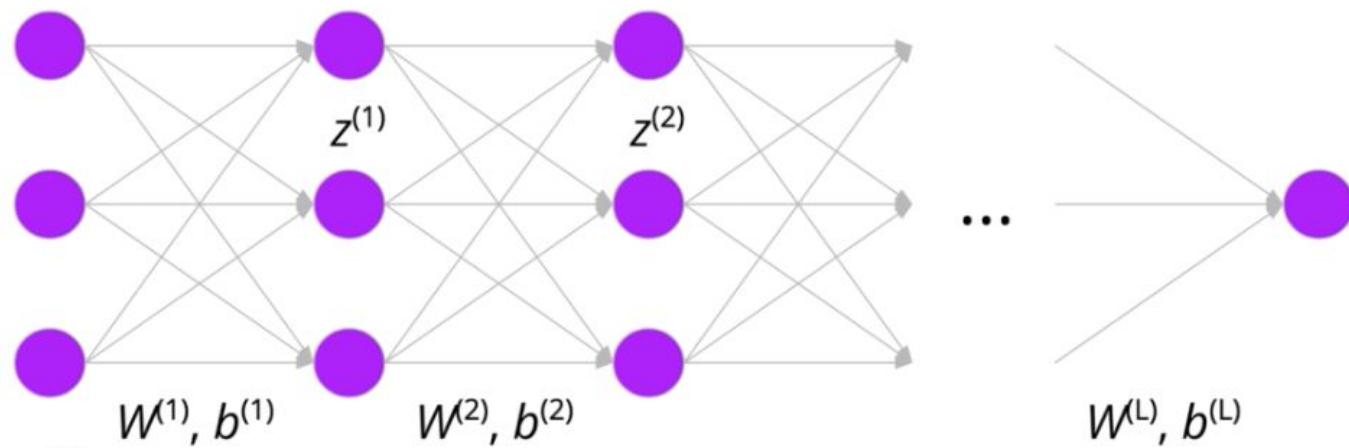
The “threshold” in this case is $-b$ (negative)

As long as $w^T x$ is $> -b$,
then $w^T x + b > 0$, and the prediction is 1



$$p(y = 1 \mid x) = \sigma(w^T x + b) = \sigma\left(\sum_{d=1}^D w_d x_d + b\right)$$

Stacking neurons



$$z^{(1)} = \sigma(W^{(1)T}x + b^{(1)})$$

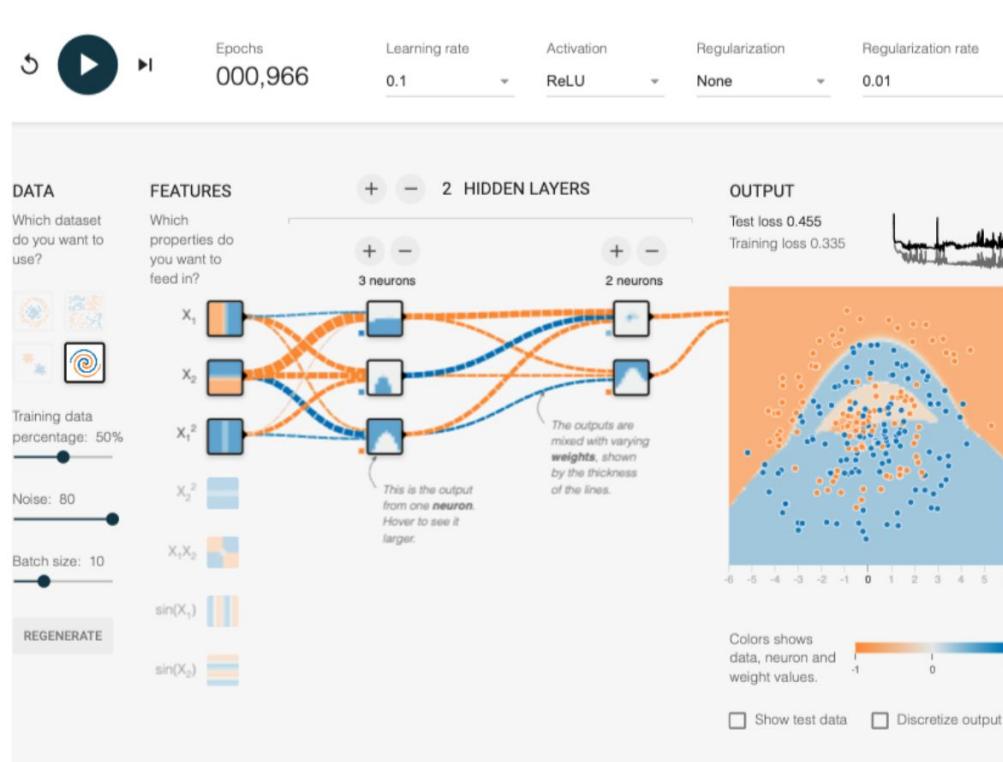
$$z^{(2)} = \sigma(W^{(2)T}z^{(1)} + b^{(2)})$$

$$z^{(3)} = \sigma(W^{(3)T}z^{(2)} + b^{(3)})$$

\dots

$$p(y = 1 \mid x) = \sigma(W^{(L)T}z^{(L-1)} + b^{(L)})$$

Fitting complex patterns through non-linear activation function and multiple layers

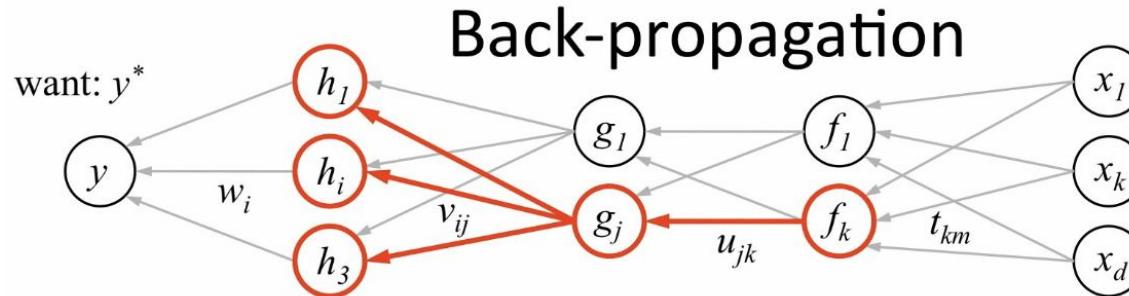


Building MLP from scratch (tensorflow and pytorch)

- <https://github.com/fastai/course22/blob/master/05-linear-model-and-neural-net-from-scratch.ipynb>
- https://colab.research.google.com/github/google/eng-edu/blob/main/ml/cc/exercises/linear_regression_with_synthetic_data.ipynb

The matrix calculus you need for deep learning

<https://explained.ai/matrix-calculus>



1. receive new observation $\mathbf{x} = [x_1 \dots x_d]$ and target y^*
2. **feed forward:** for each unit g_j in each layer 1...L
compute g_j based on units f_k from previous layer: $g_j = \sigma\left(u_{j0} + \sum_k u_{jk} f_k\right)$
3. get prediction y and error $(y - y^*)$
4. **back-propagate error:** for each unit g_j in each layer L...1

(a) compute error on g_j

$$\frac{\partial E}{\partial g_j} = \sum_i \underbrace{\sigma'(h_i)}_{\substack{\text{should } g_j \\ \text{be higher or lower?}}} \underbrace{v_{ij}}_{\substack{\text{how } h_i \text{ will} \\ \text{change as } g_j \text{ changes}}} \underbrace{\frac{\partial E}{\partial h_i}}_{\substack{\text{was } h_i \text{ too} \\ \text{high or} \\ \text{too low?}}}$$

(b) for each u_{jk} that affects g_j

(i) compute error on u_{jk}	(ii) update the weight
$\frac{\partial E}{\partial u_{jk}} = \frac{\partial E}{\partial g_j} \underbrace{\sigma'(g_j)}_{\substack{\text{do we want } g_j \text{ to} \\ \text{be higher/lower?}}} \underbrace{f_k}_{\substack{\text{how } g_j \text{ will change} \\ \text{if } u_{jk} \text{ is higher/lower?}}}$	$u_{jk} \leftarrow u_{jk} - \eta \frac{\partial E}{\partial u_{jk}}$

Multi-classification : soft-max and cross-entropy

A1	File	Edit	View	Insert	Format	Data	Tools	Help
1	entropy_example	.XLSX						
2	cat	-4.89	0.01	0.00	0	0		
3	dog	2.60	13.43	0.87	1	1		
4	plane	0.59	1.81	0.12	0	2		
5	fish	-2.07	0.13	0.01	0	3		
6	building	-4.57	0.01	0.00	0	4		
7			15.38	1.00				
8								
9								
10								
11								
12								
13								
14								
15								

e^{z_i}

$$\frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

$$H(P^* | P) = - \sum_i \underbrace{P^*(i)}_{\text{TRUE CLASS DISTRIBUTION}} \log \underbrace{P(i)}_{\text{PREDICTED CLASS DISTRIBUTION}}$$

PROPHET

Business oriented

- Built with business datasets in mind
- Model sales, CPU usage, adspend, etc.
- Do you run a business? Do you count anything? That's a time series!
- Easy to account for weekends, holidays, special events
- Easy to add custom exogenous inputs (e.g. NFL Sundays)
- Excellent choice for consultants and freelancers

Pragmatic

- Just a few lines of code and relatively fast
- Recall: ARIMA can struggle with large T
- Most things (even cross-validation) are just 1 line of code
- Very interpretable
- Reminiscent of ETS methods
- Plot components with a function call
- Seasonality at multiple scales
 - E.g. weekly / yearly
- E.g. decrease on Sunday, but increase on Black Friday

Interpretable, bayesian
for client interaction

Component of Prophet model

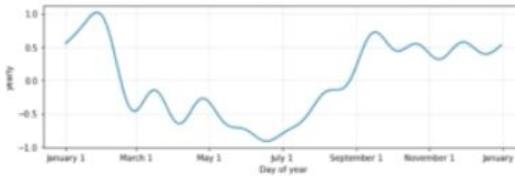
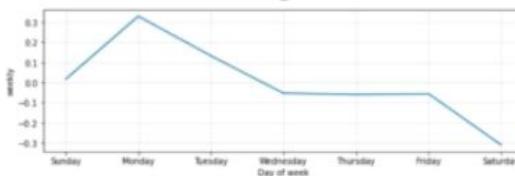
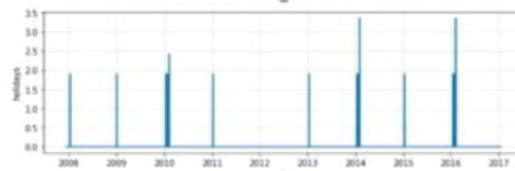
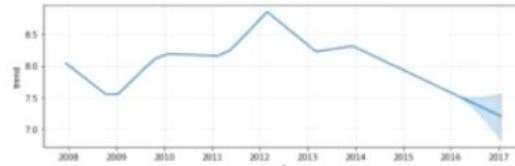
$$y(t) = g(t) + s(t) + h(t) + \varepsilon(t)$$

$g(t)$: *trend*

$s(t)$: *seasonal*

$h(t)$: *holiday*

$\varepsilon(t)$: *error*

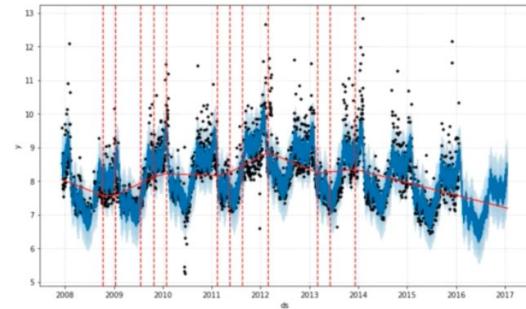


Trend $g(t)$

- Linear

slope \times time + intercept

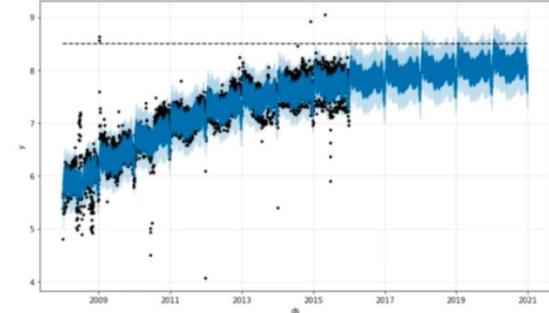
$$g(t) = (k + \mathbf{a}(t)^\top \boldsymbol{\delta})t + (m + \mathbf{a}(t)^\top \boldsymbol{\gamma})$$



- Logistic

$C(t) * \text{sigmoid}(\text{linear}(t))$

$$g(t) = \frac{C(t)}{1 + \exp(-(k + \mathbf{a}(t)^\top \boldsymbol{\delta})(t - (m + \mathbf{a}(t)^\top \boldsymbol{\gamma})))}$$



$$g(t) = \frac{C}{1 + \exp(-k(t - m))}$$

Seasonal $s(t)$

- Fourier series - don't worry if you're not an engineer / physicist
- We know that sines and cosines are periodic
- $P = 365.25$ for yearly data, $P = 7$ for weekly, when time scaled in days

$$s(t) = \sum_{n=1}^N \left(a_n \cos\left(\frac{2\pi nt}{P}\right) + b_n \sin\left(\frac{2\pi nt}{P}\right) \right)$$

Holidays $h(t)$

- One-hot encode every holiday
- Influence of each holiday governed by components of κ (additive by default)
- Can also model surrounding days, but paper doesn't specify how

$$Z(t) = [\mathbf{1}(t \in D_1), \dots, \mathbf{1}(t \in D_L)]$$

$$h(t) = Z(t)\kappa.$$

we use a prior $\kappa \sim \text{Normal}(0, \nu^2)$.

Bayesian model

- Hinted at on previous slide - all parameters have priors
- Behind the scenes, Prophet uses Stan (a Bayesian ML library)
- Allows for “human in the loop” optimization
- i.e. include your client on the modeling process (e.g. modify changepoint prior to yield the appropriate number of trend changes)

